

# **The Linux Kernel API**

## **The Linux Kernel API**

This documentation is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

For more details see the file COPYING in the source distribution of Linux.

# Table of Contents

<b>1. The Linux VFS.....</b>	<b>10</b>
The Directory Cache .....	10
d_invalidate.....	10
d_find_alias.....	10
prune_dcache .....	11
shrink_dcache_sb .....	12
have_submounts .....	12
shrink_dcache_parent .....	13
d_alloc .....	13
d_instantiate.....	14
d_alloc_root.....	15
d_lookup .....	16
d_validate.....	16
d_delete.....	17
d_rehash.....	18
d_move .....	18
__d_path .....	19
is_subdir .....	20
find_inode_number .....	21
d_drop.....	22
d_add .....	22
dget .....	23
d_unhashed .....	24
Inode Handling.....	24
__mark_inode_dirty .....	25
write_inode_now.....	25
clear_inode .....	26
invalidate_inodes .....	26
get_empty_inode .....	27
iunique .....	28
insert_inode_hash .....	29
remove_inode_hash .....	29
iput.....	30
bmap .....	30
update_atime.....	31
make_bad_inode .....	32
is_bad_inode.....	32
Registration and Superblocks.....	33

register_filesystem .....	33
unregister_filesystem .....	34
__wait_on_super .....	34
get_super .....	35
get_empty_super .....	36
<b>2. Linux Networking .....</b>	<b>37</b>
Socket Buffer Functions .....	37
skb_queue_empty .....	37
skb_get .....	37
kfree_skb .....	38
skb_cloned .....	38
skb_shared .....	39
skb_unshare .....	40
skb_peek .....	40
skb_peek_tail .....	41
skb_queue_len .....	42
__skb_queue_head .....	42
skb_queue_head .....	43
__skb_queue_tail .....	44
skb_queue_tail .....	45
__skb_dequeue .....	45
skb_dequeue .....	46
skb_insert .....	47
skb_append .....	47
skb_unlink .....	48
__skb_dequeue_tail .....	49
skb_dequeue_tail .....	49
skb_put .....	50
skb_push .....	51
skb_pull .....	51
skb_headroom .....	52
skb_tailroom .....	53
skb_reserve .....	53
skb_trim .....	54
skb_orphan .....	55
skb_queue_purge .....	55
__skb_queue_purge .....	56
dev_alloc_skb .....	57
skb_cow .....	57
skb_over_panic .....	58

skb_under_panic .....	59
alloc_skb .....	60
__kfree_skb .....	61
skb_clone .....	61
skb_copy .....	62
skb_copy_expand .....	63
Socket Filter .....	64
sk_run_filter .....	64
<b>3. Network device support.....</b>	<b>66</b>
Driver Support .....	66
init_etherdev .....	66
dev_add_pack .....	66
dev_remove_pack .....	67
__dev_get_by_name .....	68
dev_get_by_name .....	68
dev_get .....	69
__dev_get_by_index .....	70
dev_get_by_index .....	70
dev_alloc_name .....	71
dev_alloc .....	72
netdev_state_change .....	73
dev_load .....	73
dev_open .....	74
dev_close .....	75
register_netdevice_notifier .....	75
unregister_netdevice_notifier .....	76
dev_queue_xmit .....	77
netif_rx .....	77
net_call_rx_atomic .....	78
register_gifconf .....	79
netdev_set_master .....	79
dev_set_promiscuity .....	80
dev_set_allmulti .....	81
dev_ioctl .....	81
dev_new_index .....	82
register_netdevice .....	83
netdev_finish_unregister .....	84
unregister_netdevice .....	84
8390 Based Network Cards .....	85
ei_open .....	85

ei_close .....	85
ei_interrupt .....	86
ethdev_init .....	87
NS8390_init.....	87
Synchronous PPP .....	88
sppp_input .....	88
sppp_close .....	89
sppp_open .....	90
sppp_reopen.....	90
sppp_change_mtu .....	91
sppp_do_ioctl .....	92
sppp_attach .....	92
sppp_detach .....	93
<b>4. Module Loading .....</b>	<b>95</b>
request_module .....	95
<b>5. Hardware Interfaces .....</b>	<b>96</b>
Interrupt Handling .....	96
disable_irq_nosync .....	96
disable_irq .....	96
enable_irq .....	97
probe_irq_mask .....	98
MTRR Handling .....	98
mtrr_add .....	98
mtrr_del .....	100
PCI Support Library.....	100
pci_find_slot .....	101
pci_find_device.....	101
pci_find_class .....	102
pci_find_parent_resource .....	103
pci_set_power_state.....	104
pci_enable_device .....	104
MCA Architecture.....	105
MCA Device Functions .....	105
mca_find_adapter .....	105
mca_find_unused_adapter.....	106
mca_read_stored_pos.....	107
mca_read_pos .....	107
mca_write_pos .....	108
mca_set_adapter_name .....	109
mca_set_adapter_procfn .....	110

mca_is_adapter_used .....	111
mca_mark_as_used .....	111
mca_mark_as_unused .....	112
mca_get_adapter_name.....	113
mca_isadapter .....	113
mca_isenabled.....	114
MCA Bus DMA .....	114
mca_enable_dma.....	114
mca_disable_dma.....	115
mca_set_dma_addr .....	116
mca_get_dma_addr .....	116
mca_set_dma_count.....	117
mca_get_dma_residue.....	118
mca_set_dma_io .....	118
mca_set_dma_mode.....	119
<b>6. The Device File System.....</b>	<b>121</b>
devfs_register .....	121
devfs_unregister .....	122
devfs_mk_symlink .....	123
devfs_mk_dir.....	124
devfs_find_handle .....	125
devfs_get_flags.....	126
devfs_get_maj_min .....	127
devfs_get_handle_from_inode.....	128
devfs_generate_path.....	128
devfs_get_ops.....	129
devfs_set_file_size .....	130
devfs_get_info .....	131
devfs_set_info .....	131
devfs_get_parent .....	132
devfs_get_first_child .....	133
devfs_get_next_sibling.....	133
devfs_auto_unregister .....	134
devfs_get_unregister_slave .....	134
devfs_register_chrdev .....	135
devfs_register_blkdev .....	136
devfs_unregister_chrdev .....	137
devfs_unregister_blkdev .....	138
<b>7. Power Management .....</b>	<b>139</b>
pm_register.....	139

pm_unregister.....	139
pm_unregister_all.....	140
pm_send .....	141
pm_send_all .....	142
pm_find .....	142
<b>8. Miscellaneous Devices .....</b>	<b>144</b>
misc_register .....	144
misc_deregister .....	144
<b>9. Video4Linux .....</b>	<b>146</b>
video_register_device.....	146
video_unregister_device.....	146
<b>10. Sound Devices.....</b>	<b>148</b>
register_sound_special .....	148
register_sound_mixer.....	148
register_sound_midi .....	149
register_sound_dsp.....	150
register_sound_synth.....	150
unregister_sound_special .....	151
unregister_sound_mixer.....	152
unregister_sound_midi .....	152
unregister_sound_dsp.....	153
unregister_sound_synth.....	154
<b>11. 16x50 UART Driver .....</b>	<b>155</b>
register_serial .....	155
unregister_serial .....	155
<b>12. Z85230 Support Library .....</b>	<b>157</b>
z8530_interrupt .....	157
z8530_sync_open .....	157
z8530_sync_close .....	158
z8530_sync_dma_open.....	159
z8530_sync_dma_close .....	160
z8530_sync_txdma_open.....	160
z8530_sync_txdma_close .....	161
z8530_describe.....	162
z8530_init.....	162
z8530_shutdown .....	163
z8530_channel_load.....	164
z8530_null_rx .....	164
z8530_queue_xmit .....	165



z8530_get_stats .....	166
-----------------------	-----

# Chapter 1. The Linux VFS

## The Directory Cache

### d\_invalidate

**Name** `d_invalidate` — invalidate a dentry

### Synopsis

```
int d_invalidate (struct dentry * dentry);
```

### Arguments

*dentry*

dentry to invalidate

### Description

Try to invalidate the dentry if it turns out to be possible. If there are other dentries that can be reached through this one we can't delete it and we return -EBUSY. On success we return 0.

### d\_find\_alias

**Name** `d_find_alias` — grab a hashed alias of inode

## Synopsis

```
struct dentry * d_find_alias (struct inode * inode);
```

## Arguments

*inode*

inode in question

## Description

If inode has a hashed alias - acquire the reference to alias and return it. Otherwise return NULL. Notice that if inode is a directory there can be only one alias and it can be unhashed only if it has no children.

# prune\_dcache

**Name** `prune_dcache` — shrink the dcache

## Synopsis

```
void prune_dcache (int count);
```

## Arguments

*count*

number of entries to try and free

## Description

Shrink the dcache. This is done when we need more memory, or simply when we need to unmount something (at which point we need to unuse all dentries).

This function may fail to free any resources if all the dentries are in use.

## shrink\_dcache\_sb

**Name** `shrink_dcache_sb` — shrink dcache for a superblock

### Synopsis

```
void shrink_dcache_sb (struct super_block * sb);
```

### Arguments

*sb*

superblock

### Description

Shrink the dcache for the specified super block. This is used to free the dcache before unmounting a file system

## have\_submounts

**Name** `have_submounts` — check for mounts over a dentry

## Synopsis

```
int have_submounts (struct dentry * parent);
```

## Arguments

*parent*

dentry to check.

## Description

Return true if the parent or its subdirectories contain a mount point

# shrink\_dcache\_parent

**Name** `shrink_dcache_parent` — prune dcache

## Synopsis

```
void shrink_dcache_parent (struct dentry * parent);
```

## Arguments

*parent*

parent of entries to prune

## Description

Prune the dcache to remove unused children of the parent dentry.

## d\_alloc

**Name** `d_alloc` — allocate a dcache entry

### Synopsis

```
struct dentry * d_alloc (struct dentry * parent, const struct qstr * name);
```

### Arguments

*parent*

parent of entry to allocate

*name*

qstr of the name

### Description

Allocates a dentry. It returns `NULL` if there is insufficient memory available. On a success the dentry is returned. The name passed in is copied and the copy passed in may be reused after this call.

## d\_instantiate

**Name** `d_instantiate` — fill in inode information for a dentry

### Synopsis

```
void d_instantiate (struct dentry * entry, struct inode * inode);
```

## Arguments

*entry*

dentry to complete

*inode*

inode to attach to this dentry

## Description

Fill in inode information in the entry.

This turns negative dentries into productive full members of society.

NOTE! This assumes that the inode count has been incremented (or otherwise set) by the caller to indicate that it is now in use by the dcache.

## d\_alloc\_root

**Name** `d_alloc_root` — allocate root dentry

## Synopsis

```
struct dentry * d_alloc_root (struct inode * root_inode);
```

## Arguments

*root\_inode*

inode to allocate the root for

## Description

Allocate a root (“/”) dentry for the inode given. The inode is instantiated and returned. `NULL` is returned if there is insufficient memory or the inode passed is `NULL`.

## d\_lookup

**Name** `d_lookup` — search for a dentry

### Synopsis

```
struct dentry * d_lookup (struct dentry * parent, struct qstr * name);
```

### Arguments

*parent*

parent dentry

*name*

qstr of name we wish to find

### Description

Searches the children of the parent dentry for the name in question. If the dentry is found its reference count is incremented and the dentry is returned. The caller must use `d_put` to free the entry when it has finished using it. `NULL` is returned on failure.

## d\_validate

**Name** `d_validate` — verify dentry provided from insecure source



## Synopsis

```
int d_validate (struct dentry * dentry, struct dentry * dparent, unsigned int
hash, unsigned int len);
```

## Arguments

*dentry*

The dentry alleged to be valid

*dparent*

The parent dentry

*hash*

Hash of the dentry

*len*

Length of the name

## Description

An insecure source has sent us a dentry, here we verify it. This is used by ncpfs in its readdir implementation. Zero is returned in the dentry is invalid.

## NOTE

This function does `_not_` dereference the pointers before we have validated them. We can test the pointer values, but we must not actually use them until we have found a valid copy of the pointer in kernel space..

## d\_delete

**Name** `d_delete` — delete a dentry

## Synopsis

```
void d_delete (struct dentry * dentry);
```

## Arguments

*dentry*

The dentry to delete

## Description

Turn the dentry into a negative dentry if possible, otherwise remove it from the hash queues so it can be deleted later

## d\_rehash

**Name** `d_rehash` — add an entry back to the hash

## Synopsis

```
void d_rehash (struct dentry * entry);
```

## Arguments

*entry*

dentry to add to the hash

## Description

Adds a dentry to the hash according to its name.

## d\_move

**Name** `d_move` — move a dentry

### Synopsis

```
void d_move (struct dentry * dentry, struct dentry * target);
```

### Arguments

*dentry*

entry to move

*target*

new dentry

### Description

Update the dcache to reflect the move of a file name. Negative dcache entries should not be moved in this way.

## \_\_d\_path

**Name** `__d_path` — return the path of a dentry

### Synopsis

```
char * __d_path (struct dentry * dentry, struct vfsmount * vfsmnt, struct  
dentry * root, struct vfsmount * rootmnt, char * buffer, int buflen);
```

## Arguments

*dentry*

dentry to report

*vfsmnt*

– undescribed –

*root*

– undescribed –

*rootmnt*

– undescribed –

*buffer*

buffer to return value in

*buflen*

buffer length

## Description

Convert a dentry into an ASCII path name. If the entry has been deleted the string “(deleted)” is appended. Note that this is ambiguous. Returns the buffer.

“buflen” should be `PAGE_SIZE` or more.

## is\_subdir

**Name** `is_subdir` — is new dentry a subdirectory of old\_dentry

## Synopsis

```
int is_subdir (struct dentry * new_dentry, struct dentry * old_dentry);
```

## Arguments

*new\_dentry*

new dentry

*old\_dentry*

old dentry

## Description

Returns 1 if *new\_dentry* is a subdirectory of the parent (at any depth). Returns 0 otherwise.

# find\_inode\_number

**Name** `find_inode_number` — check for dentry with name

## Synopsis

```
ino_t find_inode_number (struct dentry * dir, struct qstr * name);
```

## Arguments

*dir*

directory to check

*name*

Name to find.

## Description

Check whether a dentry already exists for the given name, and return the inode number if it has an inode. Otherwise 0 is returned.

This routine is used to post-process directory listings for filesystems using synthetic inode numbers, and is necessary to keep `getcwd` working.

## d\_drop

**Name** `d_drop` — drop a dentry

### Synopsis

```
void d_drop (struct dentry * dentry);
```

### Arguments

*dentry*

dentry to drop

### Description

`d_drop` unhashes the entry from the parent dentry hashes, so that it won't be found through a VFS lookup any more. Note that this is different from deleting the dentry - `d_delete` will try to mark the dentry negative if possible, giving a successful `_negative_` lookup, while `d_drop` will just make the cache lookup fail.

`d_drop` is used mainly for stuff that wants to invalidate a dentry for some reason (NFS timeouts or autofs deletes).

## d\_add

**Name** `d_add` — add dentry to hash queues

### Synopsis

```
void d_add (struct dentry * entry, struct inode * inode);
```

### Arguments

*entry*

dentry to add

*inode*

The inode to attach to this dentry

### Description

This adds the entry to the hash queues and initializes *inode*. The entry was actually filled in earlier during `d_alloc`.

## dget

**Name** `dget` — get a reference to a dentry

### Synopsis

```
struct dentry * dget (struct dentry * dentry);
```

## Arguments

*dentry*

dentry to get a reference to

## Description

Given a dentry or `NULL` pointer increment the reference count if appropriate and return the dentry. A dentry will not be destroyed when it has references.

## d\_unhashed

**Name** `d_unhashed` — is dentry hashed

## Synopsis

```
int d_unhashed (struct dentry * dentry);
```

## Arguments

*dentry*

entry to check

## Description

Returns true if the dentry passed is not currently hashed.



## Inode Handling

### `__mark_inode_dirty`

**Name** `__mark_inode_dirty` — internal function

#### Synopsis

```
void __mark_inode_dirty (struct inode * inode);
```

#### Arguments

*inode*

inode to mark

#### Description

Mark an inode as dirty. Callers should use `mark_inode_dirty`.

### `write_inode_now`

**Name** `write_inode_now` — write an inode to disk

#### Synopsis

```
void write_inode_now (struct inode * inode);
```

## Arguments

*inode*

inode to write to disk

## Description

This function commits an inode to disk immediately if it is dirty. This is primarily needed by knfsd.

# clear\_inode

**Name** `clear_inode` — clear an inode

## Synopsis

```
void clear_inode (struct inode * inode);
```

## Arguments

*inode*

inode to clear

## Description

This is called by the filesystem to tell us that the inode is no longer useful. We just terminate it with extreme prejudice.

## invalidate\_inodes

**Name** `invalidate_inodes` — discard the inodes on a device

### Synopsis

```
int invalidate_inodes (struct super_block * sb);
```

### Arguments

*sb*

superblock

### Description

Discard all of the inodes for a given superblock. If the discard fails because there are busy inodes then a non zero value is returned. If the discard is successful all the inodes have been discarded.

## get\_empty\_inode

**Name** `get_empty_inode` — obtain an inode

### Synopsis

```
struct inode * get_empty_inode ( void );
```

### Arguments

*void*

no arguments

## Description

This is called by things like the networking layer etc that want to get an inode without any inode number, or filesystems that allocate new inodes with no pre-existing information.

On a successful return the inode pointer is returned. On a failure a NULL pointer is returned. The returned inode is not on any superblock lists.

# iunique

**Name** `iunique` — get a unique inode number

## Synopsis

```
ino_t iunique (struct super_block * sb, ino_t max_reserved);
```

## Arguments

*sb*

superblock

*max\_reserved*

highest reserved inode number

## Description

Obtain an inode number that is unique on the system for a given superblock. This is used by file systems that have no natural permanent inode numbering system. An inode number is returned that is higher than the reserved limit but unique.

## BUGS

With a large number of inodes live on the file system this function currently becomes quite slow.

## insert\_inode\_hash

**Name** `insert_inode_hash` — hash an inode

## Synopsis

```
void insert_inode_hash (struct inode * inode);
```

## Arguments

*inode*

unhashed inode

## Description

Add an inode to the inode hash for this superblock. If the inode has no superblock it is added to a separate anonymous chain.

## remove\_inode\_hash

**Name** `remove_inode_hash` — remove an inode from the hash

## Synopsis

```
void remove_inode_hash (struct inode * inode);
```

## Arguments

*inode*

inode to unhash

## Description

Remove an inode from the superblock or anonymous hash.

# iput

**Name** `iput` — put an inode

## Synopsis

```
void iput (struct inode * inode);
```

## Arguments

*inode*

inode to put

## Description

Puts an inode, dropping its usage count. If the inode use count hits zero the inode is also then freed and may be destroyed.

## bmap

**Name** `bmap` — find a block number in a file

### Synopsis

```
int bmap (struct inode * inode, int block);
```

### Arguments

*inode*

inode of file

*block*

block to find

### Description

Returns the block number on the device holding the inode that is the disk block number for the block of the file requested. That is, asked for block 4 of inode 1 the function will return the disk block relative to the disk start that holds that block of the file.

## update\_atime

**Name** `update_atime` — update the access time

### Synopsis

```
void update_atime (struct inode * inode);
```

## Arguments

*inode*

inode accessed

## Description

Update the accessed time on an inode and mark it for writeback. This function automatically handles read only file systems and media, as well as the “noatime” flag and inode specific “noatime” markers.

# make\_bad\_inode

**Name** `make_bad_inode` — mark an inode bad due to an I/O error

## Synopsis

```
void make_bad_inode (struct inode * inode);
```

## Arguments

*inode*

Inode to mark bad

## Description

When an inode cannot be read due to a media or remote network failure this function makes the inode “bad” and causes I/O operations on it to fail from this point on.



## is\_bad\_inode

**Name** `is_bad_inode` — is an inode errored

### Synopsis

```
int is_bad_inode (struct inode * inode);
```

### Arguments

*inode*

inode to test

### Description

Returns true if the inode in question has been marked as bad.

## Registration and Superblocks

## register\_filesystem

**Name** `register_filesystem` — register a new filesystem

### Synopsis

```
int register_filesystem (struct file_system_type * fs);
```

## Arguments

*fs*

the file system structure

## Description

Adds the file system passed to the list of file systems the kernel is aware of for mount and other syscalls. Returns 0 on success, or a negative errno code on an error.

The `&struct file_system_type` that is passed is linked into the kernel structures and must not be freed until the file system has been unregistered.

# unregister\_filesystem

**Name** `unregister_filesystem` — unregister a file system

## Synopsis

```
int unregister_filesystem (struct file_system_type * fs);
```

## Arguments

*fs*

filesystem to unregister

## Description

Remove a file system that was previously successfully registered with the kernel. An error is returned if the file system is not found. Zero is returned on a success.

Once this function has returned the `&struct file_system_type` structure may be freed or reused.

## **\_\_wait\_on\_super**

**Name** `__wait_on_super` — wait on a superblock

### **Synopsis**

```
void __wait_on_super (struct super_block * sb);
```

### **Arguments**

*sb*

superblock to wait on

### **Description**

Waits for a superblock to become unlocked and then returns. It does not take the lock. This is an internal function. See `wait_on_super`.

## **get\_super**

**Name** `get_super` — get the superblock of a device

### **Synopsis**

```
struct super_block * get_super (kdev_t dev);
```

### **Arguments**

*dev*

device to get the superblock for

## Description

Scans the superblock list and finds the superblock of the file system mounted on the device given. NULL is returned if no match is found.

## get\_empty\_super

**Name** `get_empty_super` — find empty superblocks

## Synopsis

```
struct super_block * get_empty_super ( void );
```

## Arguments

*void*

no arguments

## Description

Find a superblock with no device assigned. A free superblock is found and returned. If neccessary new superblocks are allocated. NULL is returned if there are insufficient resources to complete the request.

# Chapter 2. Linux Networking

## Socket Buffer Functions

### skb\_queue\_empty

**Name** `skb_queue_empty` — check if a queue is empty

### Synopsis

```
int skb_queue_empty (struct sk_buff_head * list);
```

### Arguments

*list*  
queue head

### Description

Returns true if the queue is empty, false otherwise.

### skb\_get

**Name** `skb_get` — reference buffer

### Synopsis

```
struct sk_buff * skb_get (struct sk_buff * skb);
```

## Arguments

*skb*

buffer to reference

## Description

Makes another reference to a socket buffer and returns a pointer to the buffer.

# kfree\_skb

**Name** `kfree_skb` — free an `sk_buff`

## Synopsis

```
void kfree_skb (struct sk_buff * skb);
```

## Arguments

*skb*

buffer to free

## Description

Drop a reference to the buffer and free it if the usage count has hit zero.

## skb\_cloned

**Name** `skb_cloned` — is the buffer a clone

### Synopsis

```
int skb_cloned (struct sk_buff * skb);
```

### Arguments

*skb*

buffer to check

### Description

Returns true if the buffer was generated with `skb_clone` and is one of multiple shared copies of the buffer. Cloned buffers are shared data so must not be written to under normal circumstances.

## skb\_shared

**Name** `skb_shared` — is the buffer shared

### Synopsis

```
int skb_shared (struct sk_buff * skb);
```

### Arguments

*skb*

buffer to check

## Description

Returns true if more than one person has a reference to this buffer.

# skb\_unshare

**Name** `skb_unshare` — make a copy of a shared buffer

## Synopsis

```
struct sk_buff * skb_unshare (struct sk_buff * skb, int pri);
```

## Arguments

*skb*

buffer to check

*pri*

priority for memory allocation

## Description

If the socket buffer is a clone then this function creates a new copy of the data, drops a reference count on the old copy and returns the new copy with the reference count at 1. If the buffer is not a clone the original buffer is returned. When called with a spinlock held or from interrupt state *pri* must be

`GFP_ATOMIC`

`NULL` is returned on a memory allocation failure.



## skb\_peek

**Name** `skb_peek` —

### Synopsis

```
struct sk_buff * skb_peek (struct sk_buff_head * list_);
```

### Arguments

*list\_*

list to peek at

### Description

Peek an `&sk_buff`. Unlike most other operations you **\_MUST\_** be careful with this one. A peek leaves the buffer on the list and someone else may run off with it. You must hold the appropriate locks or have a private queue to do this.

Returns `NULL` for an empty list or a pointer to the head element. The reference count is not incremented and the reference is therefore volatile. Use with caution.

## skb\_peek\_tail

**Name** `skb_peek_tail` —

### Synopsis

```
struct sk_buff * skb_peek_tail (struct sk_buff_head * list_);
```

## Arguments

*list\_*

list to peek at

## Description

Peek an `&sk_buff`. Unlike most other operations you *MUST* be careful with this one. A peek leaves the buffer on the list and someone else may run off with it. You must hold the appropriate locks or have a private queue to do this.

Returns `NULL` for an empty list or a pointer to the tail element. The reference count is not incremented and the reference is therefore volatile. Use with caution.

## skb\_queue\_len

**Name** `skb_queue_len` — get queue length

## Synopsis

```
__u32 skb_queue_len (struct sk_buff_head * list_);
```

## Arguments

*list\_*

list to measure

## Description

Return the length of an `&sk_buff` queue.

## **\_\_skb\_queue\_head**

**Name** `__skb_queue_head` — queue a buffer at the list head

### **Synopsis**

```
void __skb_queue_head (struct sk_buff_head * list, struct sk_buff * newsk);
```

### **Arguments**

*list*

list to use

*newsk*

buffer to queue

### **Description**

Queue a buffer at the start of a list. This function takes no locks and you must therefore hold required locks before calling it.

A buffer cannot be placed on two lists at the same time.

## **skb\_queue\_head**

**Name** `skb_queue_head` — queue a buffer at the list head

### **Synopsis**

```
void skb_queue_head (struct sk_buff_head * list, struct sk_buff * newsk);
```

## Arguments

*list*

list to use

*newsk*

buffer to queue

## Description

Queue a buffer at the start of the list. This function takes the list lock and can be used safely with other locking &sk\_buff functions safely.

A buffer cannot be placed on two lists at the same time.

## \_\_skb\_queue\_tail

**Name** \_\_skb\_queue\_tail — queue a buffer at the list tail

## Synopsis

```
void __skb_queue_tail (struct sk_buff_head * list, struct sk_buff * newsk);
```

## Arguments

*list*

list to use

*newsk*

buffer to queue

## Description

Queue a buffer at the end of a list. This function takes no locks and you must therefore hold required locks before calling it.

A buffer cannot be placed on two lists at the same time.

## skb\_queue\_tail

**Name** `skb_queue_tail` — queue a buffer at the list tail

## Synopsis

```
void skb_queue_tail (struct sk_buff_head * list, struct sk_buff * newsk);
```

## Arguments

*list*

list to use

*newsk*

buffer to queue

## Description

Queue a buffer at the tail of the list. This function takes the list lock and can be used safely with other locking &sk\_buff functions safely.

A buffer cannot be placed on two lists at the same time.

## **\_\_skb\_dequeue**

**Name** `__skb_dequeue` — remove from the head of the queue

### **Synopsis**

```
struct sk_buff * __skb_dequeue (struct sk_buff_head * list);
```

### **Arguments**

*list*

list to dequeue from

### **Description**

Remove the head of the list. This function does not take any locks so must be used with appropriate locks held only. The head item is returned or `NULL` if the list is empty.

## **skb\_dequeue**

**Name** `skb_dequeue` — remove from the head of the queue

### **Synopsis**

```
struct sk_buff * skb_dequeue (struct sk_buff_head * list);
```

### **Arguments**

*list*

list to dequeue from

## Description

Remove the head of the list. The list lock is taken so the function may be used safely with other locking list functions. The head item is returned or `NULL` if the list is empty.

## skb\_insert

**Name** `skb_insert` — insert a buffer

## Synopsis

```
void skb_insert (struct sk_buff * old, struct sk_buff * newsk);
```

## Arguments

*old*

buffer to insert before

*newsk*

buffer to insert

## Description

Place a packet before a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls. A buffer cannot be placed on two lists at the same time.

## skb\_append

**Name** `skb_append` — append a buffer

### Synopsis

```
void skb_append (struct sk_buff * old, struct sk_buff * newsk);
```

### Arguments

*old*

buffer to insert after

*newsk*

buffer to insert

### Description

Place a packet after a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls. A buffer cannot be placed on two lists at the same time.

## skb\_unlink

**Name** `skb_unlink` — remove a buffer from a list

### Synopsis

```
void skb_unlink (struct sk_buff * skb);
```



## Arguments

*skb*

buffer to remove

## Description

Place a packet after a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls

Works even without knowing the list it is sitting on, which can be handy at times. It also means that THE LIST MUST EXIST when you unlink. Thus a list must have its contents unlinked before it is destroyed.

## **\_\_skb\_dequeue\_tail**

**Name** `__skb_dequeue_tail` — remove from the tail of the queue

## Synopsis

```
struct sk_buff * __skb_dequeue_tail (struct sk_buff_head * list);
```

## Arguments

*list*

list to dequeue from

## Description

Remove the tail of the list. This function does not take any locks so must be used with appropriate locks held only. The tail item is returned or `NULL` if the list is empty.

## skb\_dequeue\_tail

**Name** `skb_dequeue_tail` — remove from the head of the queue

### Synopsis

```
struct sk_buff * skb_dequeue_tail (struct sk_buff_head * list);
```

### Arguments

*list*

list to dequeue from

### Description

Remove the head of the list. The list lock is taken so the function may be used safely with other locking list functions. The tail item is returned or `NULL` if the list is empty.

## skb\_put

**Name** `skb_put` — add data to a buffer

### Synopsis

```
unsigned char * skb_put (struct sk_buff * skb, unsigned int len);
```

### Arguments

*skb*

buffer to use

*len*

amount of data to add

## Description

This function extends the used data area of the buffer. If this would exceed the total buffer size the kernel will panic. A pointer to the first byte of the extra data is returned.

# skb\_push

**Name** `skb_push` — add data to the start of a buffer

## Synopsis

```
unsigned char * skb_push (struct sk_buff * skb, unsigned int len);
```

## Arguments

*skb*

buffer to use

*len*

amount of data to add

## Description

This function extends the used data area of the buffer at the buffer start. If this would exceed the total buffer headroom the kernel will panic. A pointer to the first byte of the extra data is returned.

## skb\_pull

**Name** `skb_pull` — remove data from the start of a buffer

### Synopsis

```
unsigned char * skb_pull (struct sk_buff * skb, unsigned int len);
```

### Arguments

*skb*

buffer to use

*len*

amount of data to remove

### Description

This function removes data from the start of a buffer, returning the memory to the headroom. A pointer to the next data in the buffer is returned. Once the data has been pulled future pushes will overwrite the old data.

## skb\_headroom

**Name** `skb_headroom` — bytes at buffer head

### Synopsis

```
int skb_headroom (const struct sk_buff * skb);
```

## Arguments

*skb*

buffer to check

## Description

Return the number of bytes of free space at the head of an `&sk_buff`.

# skb\_tailroom

**Name** `skb_tailroom` — bytes at buffer end

## Synopsis

```
int skb_tailroom (const struct sk_buff * skb);
```

## Arguments

*skb*

buffer to check

## Description

Return the number of bytes of free space at the tail of an `sk_buff`

## skb\_reserve

**Name** `skb_reserve` — adjust headroom

### Synopsis

```
void skb_reserve (struct sk_buff * skb, unsigned int len);
```

### Arguments

*skb*

buffer to alter

*len*

bytes to move

### Description

Increase the headroom of an empty `&sk_buff` by reducing the tail room. This is only allowed for an empty buffer.

## skb\_trim

**Name** `skb_trim` — remove end from a buffer

### Synopsis

```
void skb_trim (struct sk_buff * skb, unsigned int len);
```

## Arguments

*skb*

buffer to alter

*len*

new length

## Description

Cut the length of a buffer down by removing data from the tail. If the buffer is already under the length specified it is not modified.

# skb\_orphan

**Name** `skb_orphan` — orphan a buffer

## Synopsis

```
void skb_orphan (struct sk_buff * skb);
```

## Arguments

*skb*

buffer to orphan

## Description

If a buffer currently has an owner then we call the owner's destructor function and make the *skb* unowned. The buffer continues to exist but is no longer charged to its former owner.

## skb\_queue\_purge

**Name** `skb_queue_purge` — empty a list

### Synopsis

```
void skb_queue_purge (struct sk_buff_head * list);
```

### Arguments

*list*

list to empty

### Description

Delete all buffers on an `&sk_buff` list. Each buffer is removed from the list and one reference dropped. This function takes the list lock and is atomic with respect to other list locking functions.

## \_\_skb\_queue\_purge

**Name** `__skb_queue_purge` — empty a list

### Synopsis

```
void __skb_queue_purge (struct sk_buff_head * list);
```

### Arguments



*list*

list to empty

## Description

Delete all buffers on an &sk\_buff list. Each buffer is removed from the list and one reference dropped. This function does not take the list lock and the caller must hold the relevant locks to use it.

## dev\_alloc\_skb

**Name** dev\_alloc\_skb — allocate an skbuff for sending

## Synopsis

```
struct sk_buff * dev_alloc_skb (unsigned int length);
```

## Arguments

*length*

length to allocate

## Description

Allocate a new &sk\_buff and assign it a usage count of one. The buffer has unspecified headroom built in. Users should allocate the headroom they think they need without accounting for the built in space. The built in space is used for optimisations.

NULL is returned if there is no free memory. Although this function allocates memory it can be called from an interrupt.

## skb\_cow

**Name** `skb_cow` — copy a buffer if need be

### Synopsis

```
struct sk_buff * skb_cow (struct sk_buff * skb, unsigned int headroom);
```

### Arguments

*skb*

buffer to copy

*headroom*

needed headroom

### Description

If the buffer passed lacks sufficient headroom or is a clone then it is copied and the additional headroom made available. If there is no free memory `NULL` is returned. The new buffer is returned if a copy was made (and the old one dropped a reference). The existing buffer is returned otherwise.

This function primarily exists to avoid making two copies when making a writable copy of a buffer and then growing the headroom.

## skb\_over\_panic

**Name** `skb_over_panic` — private function

## Synopsis

```
void skb_over_panic (struct sk_buff * skb, int sz, void * here);
```

## Arguments

*skb*

buffer

*sz*

size

*here*

address

## Description

Out of line support code for `skb_put`. Not user callable.

## `skb_under_panic`

**Name** `skb_under_panic` — private function

## Synopsis

```
void skb_under_panic (struct sk_buff * skb, int sz, void * here);
```

## Arguments

*skb*  
     buffer

*sz*  
     size

*here*  
     address

## Description

Out of line support code for `skb_push`. Not user callable.

## alloc\_skb

**Name** `alloc_skb` — allocate a network buffer

## Synopsis

```
struct sk_buff * alloc_skb (unsigned int size, int gfp_mask);
```

## Arguments

*size*  
     size to allocate

*gfp\_mask*  
     allocation mask

## Description

Allocate a new `&sk_buff`. The returned buffer has no headroom and a tail room of size bytes. The object has a reference count of one. The return is the buffer. On a failure the return is `NULL`.

Buffers may only be allocated from interrupts using a *gfp\_mask* of `GFP_ATOMIC`.

## **\_\_kfree\_skb**

**Name** `__kfree_skb` — private function

### **Synopsis**

```
void __kfree_skb (struct sk_buff * skb);
```

### **Arguments**

*skb*

buffer

### **Description**

Free an `sk_buff`. Release anything attached to the buffer. Clean the state. This is an internal helper function. Users should always call `kfree_skb`

## **skb\_clone**

**Name** `skb_clone` — duplicate an `sk_buff`

## Synopsis

```
struct sk_buff * skb_clone (struct sk_buff * skb, int gfp_mask);
```

## Arguments

*skb*

buffer to clone

*gfp\_mask*

allocation priority

## Description

Duplicate an `&sk_buff`. The new one is not owned by a socket. Both copies share the same packet data but not structure. The new buffer has a reference count of 1. If the allocation fails the function returns `NULL` otherwise the new buffer is returned.

If this function is called from an interrupt `gfp_mask` must be `GFP_ATOMIC`.

## skb\_copy

**Name** `skb_copy` — copy an `sk_buff`

## Synopsis

```
struct sk_buff * skb_copy (const struct sk_buff * skb, int gfp_mask);
```

## Arguments

*skb*

buffer to copy

*gfp\_mask*

allocation priority

## Description

Make a copy of both an `&sk_buff` and its data. This is used when the caller wishes to modify the data and needs a private copy of the data to alter. Returns `NULL` on failure or the pointer to the buffer on success. The returned buffer has a reference count of 1.

You must pass `GFP_ATOMIC` as the allocation priority if this function is called from an interrupt.

# skb\_copy\_expand

**Name** `skb_copy_expand` — copy and expand `sk_buff`

## Synopsis

```
struct sk_buff * skb_copy_expand (const struct sk_buff * skb, int
newheadroom, int newtailroom, int gfp_mask);
```

## Arguments

*skb*

buffer to copy

*newheadroom*

new free bytes at head

*newtailroom*

new free bytes at tail

*gfp\_mask*

allocation priority

## Description

Make a copy of both an `&sk_buff` and its data and while doing so allocate additional space.

This is used when the caller wishes to modify the data and needs a private copy of the data to alter as well as more space for new fields. Returns `NULL` on failure or the pointer to the buffer on success. The returned buffer has a reference count of 1.

You must pass `GFP_ATOMIC` as the allocation priority if this function is called from an interrupt.

## Socket Filter

### `sk_run_filter`

**Name** `sk_run_filter` — run a filter on a socket

## Synopsis

```
int sk_run_filter (struct sk_buff * skb, struct sock_filter * filter, int
flen);
```

## Arguments

*skb*

buffer to run the filter on

*filter*

filter to apply



*flen*

length of filter

## Description

Decode and apply filter instructions to the `skb->data`. Return length to keep, 0 for none. `skb` is the data we are filtering, `filter` is the array of filter instructions, and `len` is the number of filter blocks in the array.

# Chapter 3. Network device support

## Driver Support

### init\_etherdev

**Name** `init_etherdev` — Register ethernet device

### Synopsis

```
struct net_device * init_etherdev (struct net_device * dev, int sizeof_priv);
```

### Arguments

*dev*

An ethernet device structure to be filled in, or `NULL` if a new struct should be allocated.

*sizeof\_priv*

Size of additional driver-private structure to be allocated for this ethernet device

### Description

Fill in the fields of the device structure with ethernet-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size *sizeof\_priv*. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as `dev->name`, or a new structure is made, a new name string is constructed.

## dev\_add\_pack

**Name** `dev_add_pack` — add packet handler

### Synopsis

```
void dev_add_pack (struct packet_type * pt);
```

### Arguments

*pt*

packet type declaration

### Description

Add a protocol handler to the networking stack. The passed `&packet_type` is linked into kernel lists and may not be freed until it has been removed from the kernel lists.

## dev\_remove\_pack

**Name** `dev_remove_pack` — remove packet handler

### Synopsis

```
void dev_remove_pack (struct packet_type * pt);
```

### Arguments

*pt*

packet type declaration

## Description

Remove a protocol handler that was previously added to the kernel protocol handlers by `dev_add_pack`. The passed `&packet_type` is removed from the kernel lists and can be freed or reused once this function returns.

## `__dev_get_by_name`

**Name** `__dev_get_by_name` — find a device by its name

## Synopsis

```
struct net_device * __dev_get_by_name (const char * name);
```

## Arguments

*name*

name to find

## Description

Find an interface by name. Must be called under RTNL semaphore or `dev_base_lock`. If the name is found a pointer to the device is returned. If the name is not found then `NULL` is returned. The reference counters are not incremented so the caller must be careful with locks.

## dev\_get\_by\_name

**Name** `dev_get_by_name` — find a device by its name

### Synopsis

```
struct net_device * dev_get_by_name (const char * name);
```

### Arguments

*name*

name to find

### Description

Find an interface by name. This can be called from any context and does its own locking. The returned handle has the usage count incremented and the caller must use `dev_put` to release it when it is no longer needed. `NULL` is returned if no matching device is found.

## dev\_get

**Name** `dev_get` — test if a device exists

### Synopsis

```
int dev_get (const char * name);
```

## Arguments

*name*

name to test for

## Description

Test if a name exists. Returns true if the name is found. In order to be sure the name is not allocated or removed during the test the caller must hold the rtnl semaphore.

This function primarily exists for back compatibility with older drivers.

## \_\_dev\_get\_by\_index

**Name** `__dev_get_by_index` — find a device by its ifindex

## Synopsis

```
struct net_device * __dev_get_by_index (int ifindex);
```

## Arguments

*ifindex*

index of device

## Description

Search for an interface by index. Returns `NULL` if the device is not found or a pointer to the device. The device has not had its reference counter increased so the caller must be careful about locking. The caller must hold either the RTNL semaphore or `dev_base_lock`.

## dev\_get\_by\_index

**Name** `dev_get_by_index` — find a device by its `ifindex`

### Synopsis

```
struct net_device * dev_get_by_index (int ifindex);
```

### Arguments

*ifindex*

index of device

### Description

Search for an interface by index. Returns NULL if the device is not found or a pointer to the device. The device returned has had a reference added and the pointer is safe until the user calls `dev_put` to indicate they have finished with it.

## dev\_alloc\_name

**Name** `dev_alloc_name` — allocate a name for a device

### Synopsis

```
int dev_alloc_name (struct net_device * dev, const char * name);
```

## Arguments

*dev*

device

*name*

name format string

## Description

Passed a format string - eg “%ld” it will try and find a suitable id. Not efficient for many devices, not called a lot. The caller must hold the dev\_base or rtnl lock while allocating the name and adding the device in order to avoid duplicates. Returns the number of the unit assigned or a negative errno code.

## dev\_alloc

**Name** `dev_alloc` — allocate a network device and name

## Synopsis

```
struct net_device * dev_alloc (const char * name, int * err);
```

## Arguments

*name*

name format string

*err*

error return pointer



## Description

Passed a format string, eg. “%ld”, it will allocate a network device and space for the name. `NULL` is returned if no memory is available. If the allocation succeeds then the name is assigned and the device pointer returned. `NULL` is returned if the name allocation failed. The cause of an error is returned as a negative `errno` code in the variable `err` points to.

The caller must hold the `dev_base` or `RTNL` locks when doing this in order to avoid duplicate name allocations.

## netdev\_state\_change

**Name** `netdev_state_change` — device changes state

## Synopsis

```
void netdev_state_change (struct net_device * dev);
```

## Arguments

*dev*

device to cause notification

## Description

Called to indicate a device has changed state. This function calls the notifier chains for `netdev_chain` and sends a `NEWLINK` message to the routing socket.

## dev\_load

**Name** `dev_load` — load a network module

## Synopsis

```
void dev_load (const char * name);
```

## Arguments

*name*

name of interface

## Description

If a network interface is not present and the process has suitable privileges this function loads the module. If module loading is not available in this kernel then it becomes a nop.

# dev\_open

**Name** `dev_open` — prepare an interface for use.

## Synopsis

```
int dev_open (struct net_device * dev);
```

## Arguments

*dev*

device to open

## Description

Takes a device from down to up state. The device's private open function is invoked and then the multicast lists are loaded. Finally the device is moved into the up state and a `NETDEV_UP` message is sent to the netdev notifier chain.

Calling this function on an active interface is a nop. On a failure a negative errno code is returned.

## dev\_close

**Name** `dev_close` — shutdown an interface.

### Synopsis

```
int dev_close (struct net_device * dev);
```

### Arguments

*dev*

device to shutdown

### Description

This function moves an active device into down state. A `NETDEV_GOING_DOWN` is sent to the netdev notifier chain. The device is then deactivated and finally a `NETDEV_DOWN` is sent to the notifier chain.

## register\_netdevice\_notifier

**Name** `register_netdevice_notifier` — register a network notifier block

## Synopsis

```
int register_netdevice_notifier (struct notifier_block * nb);
```

## Arguments

*nb*

notifier

## Description

Register a notifier to be called when network device events occur. The notifier passed is linked into the kernel structures and must not be reused until it has been unregistered. A negative errno code is returned on a failure.

# unregister\_netdevice\_notifier

**Name** `unregister_netdevice_notifier` — unregister a network notifier block

## Synopsis

```
int unregister_netdevice_notifier (struct notifier_block * nb);
```

## Arguments

*nb*

notifier

## Description

Unregister a notifier previously registered by `register_netdevice_notifier`. The notifier is unlinked into the kernel structures and may then be reused. A negative `errno` code is returned on a failure.

## dev\_queue\_xmit

**Name** `dev_queue_xmit` — transmit a buffer

### Synopsis

```
int dev_queue_xmit (struct sk_buff * skb);
```

### Arguments

*skb*

buffer to transmit

### Description

Queue a buffer for transmission to a network device. The caller must have set the device and priority and built the buffer before calling this function. The function can be called from an interrupt.

A negative `errno` code is returned on a failure. A success does not guarantee the frame will be transmitted as it may be dropped due to congestion or traffic shaping.

## netif\_rx

**Name** `netif_rx` — post buffer to the network code

## Synopsis

```
void netif_rx (struct sk_buff * skb);
```

## Arguments

*skb*

buffer to post

## Description

This function receives a packet from a device driver and queues it for the upper (protocol) levels to process. It always succeeds. The buffer may be dropped during processing for congestion control or by the protocol layers.

## net\_call\_rx\_atomic

**Name** `net_call_rx_atomic` —

## Synopsis

```
void net_call_rx_atomic (void (*fn) (void));
```

## Arguments

*fn*

function to call

## Description

Make a function call that is atomic with respect to the protocol layers.

## register\_gifconf

**Name** `register_gifconf` — register a SIOCGIF handler

### Synopsis

```
int register_gifconf (unsigned int family, gifconf_func_t * gifconf);
```

### Arguments

*family*

Address family

*gifconf*

Function handler

### Description

Register protocol dependent address dumping routines. The handler that is passed must not be freed or reused until it has been replaced by another handler.

## netdev\_set\_master

**Name** `netdev_set_master` — set up master/slave pair

## Synopsis

```
int netdev_set_master (struct net_device * slave, struct net_device *
master);
```

## Arguments

*slave*

slave device

*master*

new master device

## Description

Changes the master device of the slave. Pass NULL to break the bonding. The caller must hold the RTNL semaphore. On a failure a negative errno code is returned. On success the reference counts are adjusted, RTM\_NEWLINK is sent to the routing socket and the function returns zero.

# dev\_set\_promiscuity

**Name** `dev_set_promiscuity` — update promiscuity count on a device

## Synopsis

```
void dev_set_promiscuity (struct net_device * dev, int inc);
```

## Arguments

*dev*

device



*inc*

modifier

## Description

Add or remove promiscuity from a device. While the count in the device remains above zero the interface remains promiscuous. Once it hits zero the device reverts back to normal filtering operation. A negative *inc* value is used to drop promiscuity on the device.

## dev\_set\_allmulti

**Name** `dev_set_allmulti` — update allmulti count on a device

## Synopsis

```
void dev_set_allmulti (struct net_device * dev, int inc);
```

## Arguments

*dev*

device

*inc*

modifier

## Description

Add or remove reception of all multicast frames to a device. While the count in the device remains above zero the interface remains listening to all interfaces. Once it hits zero the device reverts back to normal filtering operation. A negative *inc* value is used to drop the counter when releasing a resource needing all multicasts.

## dev\_ioctl

**Name** `dev_ioctl` — network device ioctl

### Synopsis

```
int dev_ioctl (unsigned int cmd, void * arg);
```

### Arguments

*cmd*

command to issue

*arg*

pointer to a struct ifreq in user space

### Description

Issue ioctl functions to devices. This is normally called by the user space syscall interfaces but can sometimes be useful for other purposes. The return value is the return from the syscall if positive or a negative errno code on error.

## dev\_new\_index

**Name** `dev_new_index` — allocate an ifindex

### Synopsis

```
int dev_new_index ( void );
```

## Arguments

*void*

no arguments

## Description

Returns a suitable unique value for a new device interface number. The caller must hold the `rtnl` semaphore to be sure it remains unique.

# register\_netdevice

**Name** `register_netdevice` — register a network device

## Synopsis

```
int register_netdevice (struct net_device * dev);
```

## Arguments

*dev*

device to register

## Description

Take a completed network device structure and add it to the kernel interfaces. A `NETDEV_REGISTER` message is sent to the netdev notifier chain. 0 is returned on success. A negative `errno` code is returned on a failure to set up the device, or if the name is a duplicate.

## BUGS

The locking appears insufficient to guarantee two parallel registers will not get the same name.

## netdev\_finish\_unregister

**Name** `netdev_finish_unregister` — complete unregistration

### Synopsis

```
int netdev_finish_unregister (struct net_device * dev);
```

### Arguments

*dev*  
device

### Description

Destroy and free a dead device. A value of zero is returned on success.

## unregister\_netdevice

**Name** `unregister_netdevice` — remove device from the kernel

### Synopsis

```
int unregister_netdevice (struct net_device * dev);
```

## Arguments

*dev*

device

## Description

This function shuts down a device interface and removes it from the kernel tables. On success 0 is returned, on a failure a negative errno code is returned.

## 8390 Based Network Cards

### ei\_open

**Name** `ei_open` — Open/initialize the board.

## Synopsis

```
int ei_open (struct net_device * dev);
```

## Arguments

*dev*

network device to initialize

## Description

This routine goes all-out, setting everything up anew at each open, even though many of these registers should only need to be set once at boot.

## ei\_close

**Name** `ei_close` — shut down network device

### Synopsis

```
int ei_close (struct net_device * dev);
```

### Arguments

*dev*

network device to close

### Description

Opposite of `ei_open`. Only used when “`ifconfig <devname> down`” is done.

## ei\_interrupt

**Name** `ei_interrupt` — handle the interrupts from an 8390

### Synopsis

```
void ei_interrupt (int irq, void * dev_id, struct pt_regs * regs);
```

### Arguments

*irq*

interrupt number

*dev\_id*

a pointer to the net\_device

*regs*

unused

## Description

Handle the ether interface interrupts. We pull packets from the 8390 via the card specific functions and fire them at the networking stack. We also handle transmit completions and wake the transmit path if necessary. We also update the counters and do other housekeeping as needed.

## ethdev\_init

**Name** `ethdev_init` — init rest of 8390 device struct

## Synopsis

```
int ethdev_init (struct net_device * dev);
```

## Arguments

*dev*

network device structure to init

## Description

Initialize the rest of the 8390 device structure. Do NOT `__init` this, as it is used by 8390 based modular drivers too.

## NS8390\_init

**Name** `NS8390_init` — initialize 8390 hardware

### Synopsis

```
void NS8390_init (struct net_device * dev, int startp);
```

### Arguments

*dev*

network device to initialize

*startp*

boolean. non-zero value to initiate chip processing

### Description

Must be called with lock held.

## Synchronous PPP

### sppp\_input

**Name** `sppp_input` — receive and process a WAN PPP frame



## Synopsis

```
void sppp_input (struct net_device * dev, struct sk_buff * skb);
```

## Arguments

*dev*

The device it arrived on

*skb*

The buffer to process

## Description

This can be called directly by cards that do not have timing constraints but is normally called from the network layer after interrupt servicing to process frames queued via `netif_rx`.

We process the options in the card. If the frame is destined for the protocol stacks then it requeues the frame for the upper level protocol. If it is a control from it is processed and discarded here.

## sppp\_close

**Name** `sppp_close` — close down a synchronous PPP or Cisco HDLC link

## Synopsis

```
int sppp_close (struct net_device * dev);
```

## Arguments

*dev*

The network device to drop the link of

## Description

This drops the logical interface to the channel. It is not done politely as we assume we will also be dropping DTR. Any timeouts are killed.

## sppp\_open

**Name** `sppp_open` — open a synchronous PPP or Cisco HDLC link

## Synopsis

```
int sppp_open (struct net_device * dev);
```

## Arguments

*dev*

Network device to activate

## Description

Close down any existing synchronous session and commence from scratch. In the PPP case this means negotiating LCP/PCP and friends, while for Cisco HDLC we simply need to start sending keepalives

## sppp\_reopen

**Name** `sppp_reopen` — notify of physical link loss

## Synopsis

```
int sppp_reopen (struct net_device * dev);
```

## Arguments

*dev*

Device that lost the link

## Description

This function informs the synchronous protocol code that the underlying link died (for example a carrier drop on X.21)

We increment the magic numbers to ensure that if the other end failed to notice we will correctly start a new session. It happens do to the nature of telco circuits is that you can lose carrier on one endonly.

Having done this we go back to negotiating. This function may be called from an interrupt context.

## sppp\_change\_mtu

**Name** sppp\_change\_mtu — Change the link MTU

## Synopsis

```
int sppp_change_mtu (struct net_device * dev, int new_mtu);
```

## Arguments

*dev*

Device to change MTU on

*new\_mtu*

New MTU

## Description

Change the MTU on the link. This can only be called with the link down. It returns an error if the link is up or the mtu is out of range.

# sppp\_do\_ioctl

**Name** `sppp_do_ioctl` — Ioctl handler for ppp/hdlc

## Synopsis

```
int sppp_do_ioctl (struct net_device * dev, struct ifreq * ifr, int cmd);
```

## Arguments

*dev*

Device subject to ioctl

*ifr*

Interface request block from the user

*cmd*

Command that is being issued

## Description

This function handles the ioctls that may be issued by the user to control the settings of a PPP/HDLC link. It does both busy and security checks. This function is intended to be wrapped by callers who wish to add additional ioctl calls of their own.

## sppp\_attach

**Name** `sppp_attach` — attach synchronous PPP/HDLC to a device

### Synopsis

```
void sppp_attach (struct ppp_device * pd);
```

### Arguments

*pd*

PPP device to initialise

### Description

This initialises the PPP/HDLC support on an interface. At the time of calling the `dev` element must point to the network device that this interface is attached to. The interface should not yet be registered.

## sppp\_detach

**Name** `sppp_detach` — release PPP resources from a device

### Synopsis

```
void sppp_detach (struct net_device * dev);
```

### Arguments

*dev*

Network device to release

## Description

Stop and free up any PPP/HDLC resources used by this interface. This must be called before the device is freed.

# Chapter 4. Module Loading

## request\_module

**Name** `request_module` — try to load a kernel module

### Synopsis

```
int request_module (const char * module_name);
```

### Arguments

*module\_name*

Name of module

### Description

Load a module using the user mode module loader. The function returns zero on success or a negative errno code on failure. Note that a successful module load does not mean the module did not then unload and exit on an error of its own. Callers must check that the service they requested is now available not blindly invoke it.

If module auto-loading support is disabled then this function becomes a no-operation.

# Chapter 5. Hardware Interfaces

## Interrupt Handling

### **disable\_irq\_nosync**

**Name** `disable_irq_nosync` — disable an irq without waiting

#### **Synopsis**

```
void inline disable_irq_nosync (unsigned int irq);
```

#### **Arguments**

*irq*

Interrupt to disable

#### **Description**

Disable the selected interrupt line. Disables of an interrupt stack. Unlike `disable_irq`, this function does not ensure existing instances of the IRQ handler have completed before returning.

This function may be called from IRQ context.

### **disable\_irq**

**Name** `disable_irq` — disable an irq and wait for completion



## Synopsis

```
void disable_irq (unsigned int irq);
```

## Arguments

*irq*

Interrupt to disable

## Description

Disable the selected interrupt line. Disables of an interrupt stack. That is for two disables you need two enables. This function waits for any pending IRQ handlers for this interrupt to complete before returning. If you use this function while holding a resource the IRQ handler may need you will deadlock.

This function may be called - with care - from IRQ context.

## **enable\_irq**

**Name** `enable_irq` — enable interrupt handling on an `irq`

## Synopsis

```
void enable_irq (unsigned int irq);
```

## Arguments

*irq*

Interrupt to enable

## Description

Re-enables the processing of interrupts on this IRQ line providing no `disable_irq` calls are now in effect.

This function may be called from IRQ context.

## probe\_irq\_mask

**Name** `probe_irq_mask` — scan a bitmap of interrupt lines

## Synopsis

```
unsigned int probe_irq_mask (unsigned long val);
```

## Arguments

*val*

mask of interrupts to consider

## Description

Scan the ISA bus interrupt lines and return a bitmap of active interrupts. The interrupt probe logic state is then returned to its previous value.

## MTRR Handling

## mtrr\_add

**Name** `mtrr_add` — Add a memory type region

## Synopsis

```
int mtrr_add (unsigned long base, unsigned long size, unsigned int type, char
increment);
```

## Arguments

*base*

Physical base address of region

*size*

Physical size of region

*type*

Type of MTRR desired

*increment*

If this is true do usage counting on the region

## Description

Memory type region registers control the caching on newer Intel and non Intel processors. This function allows drivers to request an MTRR is added. The details and hardware specifics of each processor's implementation are hidden from the caller, but nevertheless the caller should expect to need to provide a power of two size on an equivalent power of two boundary.

If the region cannot be added either because all regions are in use or the CPU cannot support it a negative value is returned. On success the register number for this entry is returned, but should be treated as a cookie only.

On a multiprocessor machine the changes are made to all processors. This is required on x86 by the Intel processors.

The available types are

MTRR\_TYPE\_UNCACHEABLE - No caching

MTRR\_TYPE\_WRITEBACK - Write data back in bursts whenever

MTRR\_TYPE\_WRCOMB - Write data back soon but allow bursts

MTRR\_TYPE\_WRTTHROUGH - Cache reads but not writes

## BUGS

Needs a quiet flag for the cases where drivers do not mind failures and do not wish system log messages to be sent.

## mtrr\_del

**Name** `mtrr_del` — delete a memory type region

## Synopsis

```
int mtrr_del (int reg, unsigned long base, unsigned long size);
```

## Arguments

*reg*

Register returned by `mtrr_add`

*base*

Physical base address

*size*

Size of region

## Description

If register is supplied then base and size are ignored. This is how drivers should call it.

Releases an MTRR region. If the usage count drops to zero the register is freed and the region returns to default state. On success the register is returned, on failure a negative error code.

## PCI Support Library

### pci\_find\_slot

**Name** `pci_find_slot` — locate PCI device from a given PCI slot

#### Synopsis

```
struct pci_dev * pci_find_slot (unsigned int bus, unsigned int devfn);
```

#### Arguments

*bus*

number of PCI bus on which desired PCI device resides

*devfn*

number of PCI slot in which desired PCI device resides

#### Description

Given a PCI bus and slot number, the desired PCI device is located in system global list of PCI devices. If the device is found, a pointer to its data structure is returned. If no device is found, `NULL` is returned.

### pci\_find\_device

**Name** `pci_find_device` — begin or continue searching for a PCI device by vendor/device id

## Synopsis

```
struct pci_dev * pci_find_device (unsigned int vendor, unsigned int device,
const struct pci_dev * from);
```

## Arguments

*vendor*

PCI vendor id to match, or `PCI_ANY_ID` to match all vendor ids

*device*

PCI device id to match, or `PCI_ANY_ID` to match all vendor ids

*from*

Previous PCI device found in search, or `NULL` for new search.

## Description

Iterates through the list of known PCI devices. If a PCI device is found with a matching *vendor* and *device*, a pointer to its device structure is returned. Otherwise, `NULL` is returned.

A new search is initiated by passing `NULL` to the *from* argument. Otherwise if *from* is not null, searches continue from that point.

## pci\_find\_class

**Name** `pci_find_class` — begin or continue searching for a PCI device by class

## Synopsis

```
struct pci_dev * pci_find_class (unsigned int class, const struct pci_dev *
from);
```

## Arguments

*class*

search for a PCI device with this class designation

*from*

Previous PCI device found in search, or NULL for new search.

## Description

Iterates through the list of known PCI devices. If a PCI device is found with a matching *class*, a pointer to its device structure is returned. Otherwise, NULL is returned.

A new search is initiated by passing NULL to the *from* argument. Otherwise if *from* is not null, searches continue from that point.

## pci\_find\_parent\_resource

**Name** `pci_find_parent_resource` — return resource region of parent bus of given region

## Synopsis

```
struct resource * pci_find_parent_resource (const struct pci_dev * dev,
struct resource * res);
```

## Arguments

*dev*

PCI device structure contains resources to be searched

*res*

child resource record for which parent is sought

## Description

For given resource region of given device, return the resource region of parent bus the given region is contained in or where it should be allocated from.

## pci\_set\_power\_state

**Name** `pci_set_power_state` — Set power management state of a device.

## Synopsis

```
int pci_set_power_state (struct pci_dev * dev, int new_state);
```

## Arguments

*dev*

PCI device for which PM is set

*new\_state*

new power management statement (0 == D0, 3 == D3, etc.)

## Description

Set power management state of a device. For transitions from state D3 it isn't as straightforward as one could assume since many devices forget their configuration space during wakeup. Returns old power state.

## pci\_enable\_device

**Name** `pci_enable_device` — Initialize device before it's used by a driver.



## Synopsis

```
int pci_enable_device (struct pci_dev * dev);
```

## Arguments

*dev*

PCI device to be initialized

## Description

Initialize device before it's used by a driver. Ask low-level code to enable I/O and memory. Wake up the device if it was suspended. Beware, this function can fail.

## MCA Architecture

### MCA Device Functions

## **mca\_find\_adapter**

**Name** `mca_find_adapter` — scan for adapters

## Synopsis

```
int mca_find_adapter (int id, int start);
```

## Arguments

*id*

MCA identification to search for

*start*

starting slot

## Description

Search the MCA configuration for adapters matching the 16bit ID given. The first time it should be called with start as zero and then further calls made passing the return value of the previous call until MCA\_NOTFOUND is returned.

Disabled adapters are not reported.

## mca\_find\_unused\_adapter

**Name** `mca_find_unused_adapter` — scan for unused adapters

## Synopsis

```
int mca_find_unused_adapter (int id, int start);
```

## Arguments

*id*

MCA identification to search for

*start*

starting slot

## Description

Search the MCA configuration for adapters matching the 16bit ID given. The first time it should be called with start as zero and then further calls made passing the return value of the previous call until MCA\_NOTFOUND is returned.

Adapters that have been claimed by drivers and those that are disabled are not reported. This function thus allows a driver to scan for further cards when some may already be driven.

## mca\_read\_stored\_pos

**Name** `mca_read_stored_pos` — read POS register from boot data

### Synopsis

```
unsigned char mca_read_stored_pos (int slot, int reg);
```

### Arguments

*slot*

slot number to read from

*reg*

register to read from

### Description

Fetch a POS value that was stored at boot time by the kernel when it scanned the MCA space. The register value is returned. Missing or invalid registers report 0.

## mca\_read\_pos

**Name** `mca_read_pos` — read POS register from card

### Synopsis

```
unsigned char mca_read_pos (int slot, int reg);
```

### Arguments

*slot*

slot number to read from

*reg*

register to read from

### Description

Fetch a POS value directly from the hardware to obtain the current value. This is much slower than `mca_read_stored_pos` and may not be invoked from interrupt context. It handles the deep magic required for onboard devices transparently.

## mca\_write\_pos

**Name** `mca_write_pos` — read POS register from card

### Synopsis

```
void mca_write_pos (int slot, int reg, unsigned char byte);
```

## Arguments

*slot*

slot number to read from

*reg*

register to read from

*byte*

byte to write to the POS registers

## Description

Store a POS value directly from the hardware. You should not normally need to use this function and should have a very good knowledge of MCA bus before you do so. Doing this wrongly can damage the hardware.

This function may not be used from interrupt context.

Note that this is technically a Bad Thing, as IBM tech stuff says you should only set POS values through their utilities. However, some devices such as the 3c523 recommend that you write back some data to make sure the configuration is consistent. I'd say that IBM is right, but I like my drivers to work.

This function can't do checks to see if multiple devices end up with the same resources, so you might see magic smoke if someone screws up.

## mca\_set\_adapter\_name

**Name** `mca_set_adapter_name` — Set the description of the card

## Synopsis

```
void mca_set_adapter_name (int slot, char* name);
```

## Arguments

*slot*

slot to name

*name*

text string for the namen

## Description

This function sets the name reported via /proc for this adapter slot. This is for user information only. Setting a name deletes any previous name.

## mca\_set\_adapter\_procfn

**Name** `mca_set_adapter_procfn` — Set the /proc callback

## Synopsis

```
void mca_set_adapter_procfn (int slot, MCA_ProcFn procfn, void* dev);
```

## Arguments

*slot*

slot to configure

*procfn*

callback function to call for /proc

*dev*

device information passed to the callback

## Description

This sets up an information callback for `/proc/mca/slot?`. The function is called with the buffer, slot, and device pointer (or some equally informative context information, or nothing, if you prefer), and is expected to put useful information into the buffer. The adapter name, ID, and POS registers get printed before this is called though, so don't do it again.

This should be called with a `NULL procfn` when a module unregisters, thus preventing kernel crashes and other such nastiness.

## mca\_is\_adapter\_used

**Name** `mca_is_adapter_used` — check if claimed by driver

## Synopsis

```
int mca_is_adapter_used (int slot);
```

## Arguments

*slot*

slot to check

## Description

Returns 1 if the slot has been claimed by a driver

## mca\_mark\_as\_used

**Name** `mca_mark_as_used` — claim an MCA device

## Synopsis

```
int mca_mark_as_used (int slot);
```

## Arguments

*slot*

slot to claim

## FIXME

should we make this threadsafe

Claim an MCA slot for a device driver. If the slot is already taken the function returns 1, if it is not taken it is claimed and 0 is returned.

## mca\_mark\_as\_unused

**Name** `mca_mark_as_unused` — release an MCA device

## Synopsis

```
void mca_mark_as_unused (int slot);
```

## Arguments

*slot*

slot to claim

## Description



Release the slot for other drives to use.

## **mca\_get\_adapter\_name**

**Name** `mca_get_adapter_name` — get the adapter description

### **Synopsis**

```
char * mca_get_adapter_name (int slot);
```

### **Arguments**

*slot*

slot to query

### **Description**

Return the adapter description if set. If it has not been set or the slot is out range then return NULL.

## **mca\_isadapter**

**Name** `mca_isadapter` — check if the slot holds an adapter

### **Synopsis**

```
int mca_isadapter (int slot);
```

## Arguments

*slot*

slot to query

## Description

Returns zero if the slot does not hold an adapter, non zero if it does.

# mca\_isenabled

**Name** `mca_isenabled` — check if the slot holds an adapter

## Synopsis

```
int mca_isenabled (int slot);
```

## Arguments

*slot*

slot to query

## Description

Returns a non zero value if the slot holds an enabled adapter and zero for any other case.

## MCA Bus DMA

## mca\_enable\_dma

**Name** `mca_enable_dma` — channel to enable DMA on

### Synopsis

```
void mca_enable_dma (unsigned int dmanr);
```

### Arguments

*dmanr*

DMA channel

### Description

Enable the MCA bus DMA on a channel. This can be called from IRQ context.

## mca\_disable\_dma

**Name** `mca_disable_dma` — channel to disable DMA on

### Synopsis

```
void mca_disable_dma (unsigned int dmanr);
```

### Arguments

*dmanr*

DMA channel

## Description

Enable the MCA bus DMA on a channel. This can be called from IRQ context.

## mca\_set\_dma\_addr

**Name** `mca_set_dma_addr` — load a 24bit DMA address

## Synopsis

```
void mca_set_dma_addr (unsigned int dmanr, unsigned int a);
```

## Arguments

*dmanr*

DMA channel

*a*

24bit bus address

## Description

Load the address register in the DMA controller. This has a 24bit limitation (16Mb).

## mca\_get\_dma\_addr

**Name** `mca_get_dma_addr` — load a 24bit DMA address

### Synopsis

```
unsigned int mca_get_dma_addr (unsigned int dmanr);
```

### Arguments

*dmanr*

DMA channel

### Description

Read the address register in the DMA controller. This has a 24bit limitation (16Mb). The return is a bus address.

## mca\_set\_dma\_count

**Name** `mca_set_dma_count` — load a 16bit transfer count

### Synopsis

```
void mca_set_dma_count (unsigned int dmanr, unsigned int count);
```

### Arguments

*dmanr*

DMA channel

*count*

count

## Description

Set the DMA count for this channel. This can be up to 64Kbytes. Setting a count of zero will not do what you expect.

## mca\_get\_dma\_residue

**Name** `mca_get_dma_residue` — get the remaining bytes to transfer

## Synopsis

```
unsigned int mca_get_dma_residue (unsigned int dmanr);
```

## Arguments

*dmanr*

DMA channel

## Description

This function returns the number of bytes left to transfer on this DMA channel.

## mca\_set\_dma\_io

**Name** `mca_set_dma_io` — set the port for an I/O transfer

### Synopsis

```
void mca_set_dma_io (unsigned int dmanr, unsigned int io_addr);
```

### Arguments

*dmanr*

DMA channel

*io\_addr*

an I/O port number

### Description

Unlike the ISA bus DMA controllers the DMA on MCA bus can transfer with an I/O port target.

## mca\_set\_dma\_mode

**Name** `mca_set_dma_mode` — set the DMA mode

### Synopsis

```
void mca_set_dma_mode (unsigned int dmanr, unsigned int mode);
```

## Arguments

*dmanr*

DMA channel

*mode*

mode to set

## Description

The DMA controller supports several modes. The mode values you can

### set are

MCA\_DMA\_MODE\_READ when reading from the DMA device.

MCA\_DMA\_MODE\_WRITE to writing to the DMA device.

MCA\_DMA\_MODE\_IO to do DMA to or from an I/O port.

MCA\_DMA\_MODE\_16 to do 16bit transfers.



# Chapter 6. The Device File System

## devfs\_register

**Name** `devfs_register` — Register a device entry.

### Synopsis

```
devfs_handle_t devfs_register (devfs_handle_t dir, const char * name,  
unsigned int namelen, unsigned int flags, unsigned int major, unsigned int  
minor, umode_t mode, uid_t uid, gid_t gid, void * ops, void * info);
```

### Arguments

*dir*

The handle to the parent devfs directory entry. If this is `NULL` the new name is relative to the root of the devfs.

*name*

The name of the entry.

*namelen*

The number of characters in *name*, not including a `NULL` terminator. If this is 0, then *name* must be `NULL`-terminated and the length is computed internally.

*flags*

A set of bitwise-ORed flags (`DEVFS_FL_*`).

*major*

The major number. Not needed for regular files.

*minor*

The minor number. Not needed for regular files.

*mode*

The default file mode.

*uid*

The default UID of the file.

*gid*

– undescribed –

*ops*The `&file_operations` or `&block_device_operations` structure. This must not be externally deallocated.*info*An arbitrary pointer which will be written to the *private\_data* field of the `&file` structure passed to the device driver. You can set this to whatever you like, and change it once the file is opened (the next file opened will not see this change).

## Description

Returns a handle which may later be used in a call to `devfs_unregister`. On failure `NULL` is returned.

## devfs\_unregister

**Name** `devfs_unregister` — Unregister a device entry.

## Synopsis

```
void devfs_unregister (devfs_handle_t de);
```

## Arguments

*de*

– undescribed –

**de**

A handle previously created by `devfs_register` or returned from `devfs_find_handle`. If this is `NULL` the routine does nothing.

## **devfs\_mk\_symlink**

**Name** `devfs_mk_symlink` —

### **Synopsis**

```
int devfs_mk_symlink (devfs_handle_t dir, const char * name, unsigned int
namelen, unsigned int flags, const char * link, unsigned int linklength,
devfs_handle_t * handle, void * info);
```

### **Arguments**

*dir*

The handle to the parent devfs directory entry. If this is `NULL` the new name is relative to the root of the devfs.

*name*

The name of the entry.

*namelen*

The number of characters in *name*, not including a `NULL` terminator. If this is 0, then *name* must be `NULL`-terminated and the length is computed internally.

*flags*

A set of bitwise-ORed flags (DEVFS\_FL\_\*).

*link*

The destination name.

*linklength*

The number of characters in *link*, not including a NULL terminator. If this is 0, then *link* must be NULL-terminated and the length is computed internally.

*handle*

The handle to the symlink entry is written here. This may be NULL.

*info*

An arbitrary pointer which will be associated with the entry.

## Description

Returns 0 on success, else a negative error code is returned.

## devfs\_mk\_dir

**Name** `devfs_mk_dir` — Create a directory in the devfs namespace.

## Synopsis

```
devfs_handle_t devfs_mk_dir (devfs_handle_t dir, const char * name, unsigned
int namelen, void * info);
```

## Arguments

*dir*

The handle to the parent devfs directory entry. If this is `NULL` the new name is relative to the root of the devfs.

*name*

The name of the entry.

*namelen*

The number of characters in *name*, not including a `NULL` terminator. If this is 0, then *name* must be `NULL`-terminated and the length is computed internally.

*info*

An arbitrary pointer which will be associated with the entry.

## Description

Use of this function is optional. The `devfs_register` function will automatically create intermediate directories as needed. This function is provided for efficiency reasons, as it provides a handle to a directory. Returns a handle which may later be used in a call to `devfs_unregister`. On failure `NULL` is returned.

## devfs\_find\_handle

**Name** `devfs_find_handle` — Find the handle of a devfs entry.

## Synopsis

```
devfs_handle_t devfs_find_handle (devfs_handle_t dir, const char * name,
unsigned int namelen, unsigned int major, unsigned int minor, char type, int
traverse_symlinks);
```

## Arguments

*dir*

The handle to the parent devfs directory entry. If this is `NULL` the name is relative to the root of the devfs.

*name*

The name of the entry.

*namelen*

The number of characters in *name*, not including a `NULL` terminator. If this is 0, then *name* must be `NULL`-terminated and the length is computed internally.

*major*

The major number. This is used if *name* is `NULL`.

*minor*

The minor number. This is used if *name* is `NULL`.

*type*

The type of special file to search for. This may be either `DEVFS_SPECIAL_CHR` or `DEVFS_SPECIAL_BLK`.

*traverse\_symlinks*

If `TRUE` then symlink entries in the devfs namespace are traversed. Symlinks pointing out of the devfs namespace will cause a failure. Symlink traversal consumes stack space.

## Description

Returns a handle which may later be used in a call to `devfs_unregister`, `devfs_get_flags`, or `devfs_set_flags`. On failure `NULL` is returned.

## devfs\_get\_flags

**Name** `devfs_get_flags` — Get the flags for a devfs entry.

## Synopsis

```
int devfs_get_flags (devfs_handle_t de, unsigned int * flags);
```

## Arguments

*de*

The handle to the device entry.

*flags*

The flags are written here.

## Description

Returns 0 on success, else a negative error code.

# devfs\_get\_maj\_min

**Name** `devfs_get_maj_min` — Get the major and minor numbers for a devfs entry.

## Synopsis

```
int devfs_get_maj_min (devfs_handle_t de, unsigned int * major, unsigned int * minor);
```

## Arguments

*de*

The handle to the device entry.

*major*

The major number is written here. This may be `NULL`.

*minor*

The minor number is written here. This may be `NULL`.

## Description

Returns 0 on success, else a negative error code.

# devfs\_get\_handle\_from\_inode

**Name** `devfs_get_handle_from_inode` — Get the devfs handle for a VFS inode.

## Synopsis

```
devfs_handle_t devfs_get_handle_from_inode (struct inode * inode);
```

## Arguments

*inode*

The VFS inode.

## Description

Returns the devfs handle on success, else `NULL`.



## devfs\_generate\_path

**Name** `devfs_generate_path` — Generate a pathname for an entry, relative to the devfs root.

### Synopsis

```
int devfs_generate_path (devfs_handle_t de, char * path, int buflen);
```

### Arguments

*de*

The devfs entry.

*path*

The buffer to write the pathname to. The pathname and '\0' terminator will be written at the end of the buffer.

*buflen*

The length of the buffer.

### Description

Returns the offset in the buffer where the pathname starts on success, else a negative error code.

## devfs\_get\_ops

**Name** `devfs_get_ops` — Get the device operations for a devfs entry.

## Synopsis

```
void * devfs_get_ops (devfs_handle_t de);
```

## Arguments

*de*

The handle to the device entry.

## Description

Returns a pointer to the device operations on success, else NULL.

## devfs\_set\_file\_size

**Name** `devfs_set_file_size` — Set the file size for a devfs regular file.

## Synopsis

```
int devfs_set_file_size (devfs_handle_t de, unsigned long size);
```

## Arguments

*de*

– undescribed –

*size*

– undescribed –

**de**

The handle to the device entry.

**size**

The new file size.

Returns 0 on success, else a negative error code.

**devfs\_get\_info**

**Name** `devfs_get_info` — Get the info pointer written to `private_data` of `@de` upon open.

**Synopsis**

```
void * devfs_get_info (devfs_handle_t de);
```

**Arguments**

*de*

The handle to the device entry.

**Description**

Returns the info pointer.

**devfs\_set\_info**

**Name** `devfs_set_info` — Set the info pointer written to `private_data` upon open.

## Synopsis

```
int devfs_set_info (devfs_handle_t de, void * info);
```

## Arguments

*de*

The handle to the device entry.

*info*

– undescribed –

## Description

Returns 0 on success, else a negative error code.

## devfs\_get\_parent

**Name** `devfs_get_parent` — Get the parent device entry.

## Synopsis

```
devfs_handle_t devfs_get_parent (devfs_handle_t de);
```

## Arguments

*de*

The handle to the device entry.

## Description

Returns the parent device entry if it exists, else NULL.

## devfs\_get\_first\_child

**Name** `devfs_get_first_child` — Get the first leaf node in a directory.

## Synopsis

```
devfs_handle_t devfs_get_first_child (devfs_handle_t de);
```

## Arguments

*de*

The handle to the device entry.

## Description

Returns the leaf node device entry if it exists, else NULL.

## devfs\_get\_next\_sibling

**Name** `devfs_get_next_sibling` — Get the next sibling leaf node. for a device entry.

## Synopsis

```
devfs_handle_t devfs_get_next_sibling (devfs_handle_t de);
```

## Arguments

*de*

The handle to the device entry.

## Description

Returns the leaf node device entry if it exists, else `NULL`.

# devfs\_auto\_unregister

**Name** `devfs_auto_unregister` — Configure a devfs entry to be automatically unregistered.

## Synopsis

```
void devfs_auto_unregister (devfs_handle_t master, devfs_handle_t slave);
```

## Arguments

*master*

The master devfs entry. Only one slave may be registered.

*slave*

The devfs entry which will be automatically unregistered when the master entry is unregistered. It is illegal to call `devfs_unregister` on this entry.

## devfs\_get\_unregister\_slave

**Name** `devfs_get_unregister_slave` — Get the slave entry which will be automatically unregistered.

### Synopsis

```
devfs_handle_t devfs_get_unregister_slave (devfs_handle_t master);
```

### Arguments

*master*

The master devfs entry.

### Description

Returns the slave which will be unregistered when *master* is unregistered.

## devfs\_register\_chrdev

**Name** `devfs_register_chrdev` — Optionally register a conventional character driver.

### Synopsis

```
int devfs_register_chrdev (unsigned int major, const char * name, struct
file_operations * fops);
```

## Arguments

*major*

The major number for the driver.

*name*

The name of the driver (as seen in /proc/devices).

*fops*

The &file\_operations structure pointer.

## Description

This function will register a character driver provided the “devfs=only” option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

## devfs\_register\_blkdev

**Name** `devfs_register_blkdev` — Optionally register a conventional block driver.

## Synopsis

```
int devfs_register_blkdev (unsigned int major, const char * name, struct
block_device_operations * bdops);
```

## Arguments

*major*

The major number for the driver.

*name*

The name of the driver (as seen in /proc/devices).



*bdops*

The `&block_device_operations` structure pointer.

## Description

This function will register a block driver provided the “devfs=only” option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

# devfs\_unregister\_chrdev

**Name** `devfs_unregister_chrdev` — Optionally unregister a conventional character driver.

## Synopsis

```
int devfs_unregister_chrdev (unsigned int major, const char * name);
```

## Arguments

*major*

– undescribed –

*name*

– undescribed –

## major

The major number for the driver.

## name

The name of the driver (as seen in `/proc/devices`).

This function will unregister a character driver provided the “devfs=only” option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

## devfs\_unregister\_blkdev

**Name** `devfs_unregister_blkdev` — Optionally unregister a conventional block driver.

### Synopsis

```
int devfs_unregister_blkdev (unsigned int major, const char * name);
```

### Arguments

*major*

The major number for the driver.

*name*

The name of the driver (as seen in `/proc/devices`).

### Description

This function will unregister a block driver provided the “devfs=only” option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

# Chapter 7. Power Management

## pm\_register

**Name** `pm_register` — register a device with power management

### Synopsis

```
struct pm_dev * pm_register (pm_dev_t type, unsigned long id, pm_callback  
callback);
```

### Arguments

*type*

device type

*id*

device ID

*callback*

callback function

### Description

Add a device to the list of devices that wish to be notified about power management events. A `&pm_dev` structure is returned on success, on failure the return is `NULL`.

## pm\_unregister

**Name** `pm_unregister` — unregister a device with power management

### Synopsis

```
void pm_unregister (struct pm_dev * dev);
```

### Arguments

*dev*

device to unregister

### Description

Remove a device from the power management notification lists. The dev passed must be a handle previously returned by `pm_register`.

## pm\_unregister\_all

**Name** `pm_unregister_all` — unregister all devices with matching callback

### Synopsis

```
void pm_unregister_all (pm_callback callback);
```

### Arguments

*callback*

callback function pointer

## Description

Unregister every device that would call the callback passed. This is primarily meant as a helper function for loadable modules. It enables a module to give up all its managed devices without keeping its own private list.

## pm\_send

**Name** `pm_send` — send request to a single device

## Synopsis

```
int pm_send (struct pm_dev * dev, pm_request_t rqst, void * data);
```

## Arguments

*dev*

device to send to

*rqst*

power management request

*data*

data for the callback

## Description

Issue a power management request to a given device. The `PM_SUSPEND` and `PM_RESUME` events are handled specially. The data field must hold the intended next state. No call is made if the state matches.

## BUGS

what stops two power management requests occuring in parallel and conflicting.

## pm\_send\_all

**Name** `pm_send_all` — send request to all managed devices

## Synopsis

```
int pm_send_all (pm_request_t rqst, void * data);
```

## Arguments

*rqst*

power management request

*data*

data for the callback

## Description

Issue a power management request to a all devices. The `PM_SUSPEND` events are handled specially. Any device is permitted to fail a suspend by returning a non zero (error) value from its callback function. If any device vetoes a suspend request then all other devices that have suspended during the processing of this request are restored to their previous state.

Zero is returned on success. If a suspend fails then the status from the device that vetoes the suspend is returned.

## BUGS

what stops two power management requests occuring in parallel and conflicting.

## pm\_find

**Name** `pm_find` — find a device

### Synopsis

```
struct pm_dev * pm_find (pm_dev_t type, struct pm_dev * from);
```

### Arguments

*type*

type of device

*from*

where to start looking

### Description

Scan the power management list for devices of a specific type. The return value for a matching device may be passed to further calls to this function to find further matches. A `NULL` indicates the end of the list.

To search from the beginning pass `NULL` as the *from* value.

## Chapter 8. Miscellaneous Devices

### **misc\_register**

**Name** `misc_register` — register a miscellaneous device

#### **Synopsis**

```
int misc_register (struct miscdevice * misc);
```

#### **Arguments**

*misc*

device structure

#### **Description**

Register a miscellaneous device with the kernel. If the minor number is set to `MISC_DYNAMIC_MINOR` a minor number is assigned and placed in the minor field of the structure. For other cases the minor number requested is used.

The structure passed is linked into the kernel and may not be destroyed until it has been unregistered.

A zero is returned on success and a negative `errno` code for failure.

### **misc\_deregister**

**Name** `misc_deregister` — unregister a miscellaneous device



## Synopsis

```
int misc_deregister (struct miscdevice * misc);
```

## Arguments

*misc*

device to unregister

## Description

Unregister a miscellaneous device that was previously successfully registered with `misc_register`. Success is indicated by a zero return, a negative `errno` code indicates an error.

# Chapter 9. Video4Linux

## video\_register\_device

**Name** `video_register_device` — register video4linux devices

### Synopsis

```
int video_register_device (struct video_device * vfd, int type);
```

### Arguments

*vfd*

video device structure we want to register

*type*

type of device to register

### FIXME

needs a semaphore on 2.3.x

The registration code assigns minor numbers based on the type requested. -ENFILE is returned in all the device slots for this category are full. If not then the minor field is set and the driver initialize function is called (if non NULL).

Zero is returned on success.

Valid types are

VFL\_TYPE\_GRABBER - A frame grabber

VFL\_TYPE\_VTX - A teletext device

VFL\_TYPE\_VBI - Vertical blank data (undecoded)

VFL\_TYPE\_RADIO - A radio card

## video\_unregister\_device

**Name** `video_unregister_device` — unregister a video4linux device

### Synopsis

```
void video_unregister_device (struct video_device * vfd);
```

### Arguments

*vfd*

the device to unregister

### Description

This unregisters the passed device and deassigns the minor number. Future open calls will be met with errors.

# Chapter 10. Sound Devices

## register\_sound\_special

**Name** `register_sound_special` — register a special sound node

### Synopsis

```
int register_sound_special (struct file_operations * fops, int unit);
```

### Arguments

*fops*

File operations for the driver

*unit*

Unit number to allocate

### Description

Allocate a special sound device by minor number from the sound subsystem. The allocated number is returned on succes. On failure a negative error code is returned.

## register\_sound\_mixer

**Name** `register_sound_mixer` — register a mixer device

## Synopsis

```
int register_sound_mixer (struct file_operations * fops, int dev);
```

## Arguments

*fops*

File operations for the driver

*dev*

Unit number to allocate

## Description

Allocate a mixer device. Unit is the number of the mixer requested. Pass -1 to request the next free mixer unit. On success the allocated number is returned, on failure a negative error code is returned.

# register\_sound\_midi

**Name** `register_sound_midi` — register a midi device

## Synopsis

```
int register_sound_midi (struct file_operations * fops, int dev);
```

## Arguments

*fops*

File operations for the driver

*dev*

Unit number to allocate

## Description

Allocate a midi device. Unit is the number of the midi device requested. Pass -1 to request the next free midi unit. On success the allocated number is returned, on failure a negative error code is returned.

# register\_sound\_dsp

**Name** `register_sound_dsp` — register a DSP device

## Synopsis

```
int register_sound_dsp (struct file_operations * fops, int dev);
```

## Arguments

*fops*

File operations for the driver

*dev*

Unit number to allocate

## Description

Allocate a DSP device. Unit is the number of the DSP requested. Pass -1 to request the next free DSP unit. On success the allocated number is returned, on failure a negative error code is returned.

This function allocates both the audio and dsp device entries together and will always allocate them as a matching pair - eg dsp3/audio3

## register\_sound\_synth

**Name** `register_sound_synth` — register a synth device

### Synopsis

```
int register_sound_synth (struct file_operations * fops, int dev);
```

### Arguments

*fops*

File operations for the driver

*dev*

Unit number to allocate

### Description

Allocate a synth device. Unit is the number of the synth device requested. Pass -1 to request the next free synth unit. On success the allocated number is returned, on failure a negative error code is returned.

## unregister\_sound\_special

**Name** `unregister_sound_special` — unregister a special sound device

### Synopsis

```
void unregister_sound_special (int unit);
```

## Arguments

*unit*

unit number to allocate

## Description

Release a sound device that was allocated with `register_sound_special`. The unit passed is the return value from the register function.

# unregister\_sound\_mixer

**Name** `unregister_sound_mixer` — unregister a mixer

## Synopsis

```
void unregister_sound_mixer (int unit);
```

## Arguments

*unit*

unit number to allocate

## Description

Release a sound device that was allocated with `register_sound_mixer`. The unit passed is the return value from the register function.



## unregister\_sound\_midi

**Name** `unregister_sound_midi` — unregister a midi device

### Synopsis

```
void unregister_sound_midi (int unit);
```

### Arguments

*unit*

unit number to allocate

### Description

Release a sound device that was allocated with `register_sound_midi`. The unit passed is the return value from the register function.

## unregister\_sound\_dsp

**Name** `unregister_sound_dsp` — unregister a DSP device

### Synopsis

```
void unregister_sound_dsp (int unit);
```

### Arguments

*unit*

unit number to allocate

## Description

Release a sound device that was allocated with `register_sound_dsp`. The unit passed is the return value from the register function.

Both of the allocated units are released together automatically.

# unregister\_sound\_synth

**Name** `unregister_sound_synth` — unregister a synth device

## Synopsis

```
void unregister_sound_synth (int unit);
```

## Arguments

*unit*

unit number to allocate

## Description

Release a sound device that was allocated with `register_sound_synth`. The unit passed is the return value from the register function.

# Chapter 11. 16x50 UART Driver

## register\_serial

**Name** `register_serial` — configure a 16x50 serial port at runtime

### Synopsis

```
int register_serial (struct serial_struct * req);
```

### Arguments

*req*

request structure

### Description

Configure the serial port specified by the request. If the port exists and is in use an error is returned. If the port is not currently in the table it is added.

The port is then probed and if necessary the IRQ is autodetected. If this fails an error is returned.

On success the port is ready to use and the line number is returned.

## unregister\_serial

**Name** `unregister_serial` — deconfigure a 16x50 serial port

## Synopsis

```
void unregister_serial (int line);
```

## Arguments

*line*

line to deconfigure

## Description

The port specified is deconfigured and its resources are freed. Any user of the port is disconnected as if carrier was dropped. Line is the port number returned by `register_serial`.

# Chapter 12. Z85230 Support Library

## z8530\_interrupt

**Name** `z8530_interrupt` — Handle an interrupt from a Z8530

### Synopsis

```
void z8530_interrupt (int irq, void * dev_id, struct pt_regs * regs);
```

### Arguments

*irq*

Interrupt number

*dev\_id*

The Z8530 device that is interrupting.

*regs*

unused

### Description

A Z85[2]30 device has stuck its hand in the air for attention. We scan both the channels on the chip for events and then call the channel specific call backs for each channel that has events. We have to use callback functions because the two channels can be in different modes.

## z8530\_sync\_open

**Name** `z8530_sync_open` — Open a Z8530 channel for PIO

### Synopsis

```
int z8530_sync_open (struct net_device * dev, struct z8530_channel * c);
```

### Arguments

*dev*

The network interface we are using

*c*

The Z8530 channel to open in synchronous PIO mode

### Description

Switch a Z8530 into synchronous mode without DMA assist. We raise the RTS/DTR and commence network operation.

## z8530\_sync\_close

**Name** `z8530_sync_close` — Close a PIO Z8530 channel

### Synopsis

```
int z8530_sync_close (struct net_device * dev, struct z8530_channel * c);
```

## Arguments

*dev*

Network device to close

*c*

Z8530 channel to disassociate and move to idle

## Description

Close down a Z8530 interface and switch its interrupt handlers to discard future events.

## z8530\_sync\_dma\_open

**Name** `z8530_sync_dma_open` — Open a Z8530 for DMA I/O

## Synopsis

```
int z8530_sync_dma_open (struct net_device * dev, struct z8530_channel * c);
```

## Arguments

*dev*

The network device to attach

*c*

The Z8530 channel to configure in sync DMA mode.

## Description

Set up a Z85x30 device for synchronous DMA in both directions. Two ISA DMA channels must be available for this to work. We assume ISA DMA driven I/O and PC limits on access.

## **z8530\_sync\_dma\_close**

**Name** `z8530_sync_dma_close` — Close down DMA I/O

### **Synopsis**

```
int z8530_sync_dma_close (struct net_device * dev, struct z8530_channel * c);
```

### **Arguments**

*dev*

Network device to detach

*c*

Z8530 channel to move into discard mode

### **Description**

Shut down a DMA mode synchronous interface. Halt the DMA, and free the buffers.

## **z8530\_sync\_txdma\_open**

**Name** `z8530_sync_txdma_open` — Open a Z8530 for TX driven DMA



## Synopsis

```
int z8530_sync_txdma_open (struct net_device * dev, struct z8530_channel *
c);
```

## Arguments

*dev*

The network device to attach

*c*

The Z8530 channel to configure in sync DMA mode.

## Description

Set up a Z85x30 device for synchronous DMA transmission. One ISA DMA channel must be available for this to work. The receive side is run in PIO mode, but then it has the bigger FIFO.

## z8530\_sync\_txdma\_close

**Name** z8530\_sync\_txdma\_close — Close down a TX driven DMA channel

## Synopsis

```
int z8530_sync_txdma_close (struct net_device * dev, struct z8530_channel *
c);
```

## Arguments

*dev*

Network device to detach

*c*

Z8530 channel to move into discard mode

## Description

Shut down a DMA/PIO split mode synchronous interface. Halt the DMA, and free the buffers.

## z8530\_describe

**Name** `z8530_describe` — Uniformly describe a Z8530 port

## Synopsis

```
void z8530_describe (struct z8530_dev * dev, char * mapping, unsigned long
io);
```

## Arguments

*dev*

Z8530 device to describe

*mapping*

string holding mapping type (eg “I/O” or “Mem”)

*io*

the port value in question

## Description

Describe a Z8530 in a standard format. We must pass the I/O as the port offset isn't predictable. The main reason for this function is to try and get a common format of report.

## z8530\_init

**Name** `z8530_init` — Initialise a Z8530 device

### Synopsis

```
int z8530_init (struct z8530_dev * dev);
```

### Arguments

*dev*

Z8530 device to initialise.

### Description

Configure up a Z8530/Z85C30 or Z85230 chip. We check the device is present, identify the type and then program it to hopefully keep quite and behave. This matters a lot, a Z8530 in the wrong state will sometimes get into stupid modes generating 10Khz interrupt streams and the like.

We set the interrupt handler up to discard any events, in case we get them during reset or setp.

Return 0 for success, or a negative value indicating the problem in errno form.

## z8530\_shutdown

**Name** `z8530_shutdown` — Shutdown a Z8530 device

### Synopsis

```
int z8530_shutdown (struct z8530_dev * dev);
```

## Arguments

*dev*

The Z8530 chip to shutdown

## Description

We set the interrupt handlers to silence any interrupts. We then reset the chip and wait 100uS to be sure the reset completed. Just in case the caller then tries to do stuff.

## z8530\_channel\_load

**Name** z8530\_channel\_load — Load channel data

## Synopsis

```
int z8530_channel_load (struct z8530_channel * c, u8 * rtable);
```

## Arguments

*c*

Z8530 channel to configure

*rtable*

table of register, value pairs

## FIXME

ioctl to allow user uploaded tables

Load a Z8530 channel up from the system data. We use +16 to indicate the “prime” registers. The value 255 terminates the table.

## z8530\_null\_rx

**Name** z8530\_null\_rx — Discard a packet

### Synopsis

```
void z8530_null_rx (struct z8530_channel * c, struct sk_buff * skb);
```

### Arguments

*c*

The channel the packet arrived on

*skb*

The buffer

### Description

We point the receive handler at this function when idle. Instead of syncppp processing the frames we get to throw them away.

## z8530\_queue\_xmit

**Name** z8530\_queue\_xmit — Queue a packet

### Synopsis

```
int z8530_queue_xmit (struct z8530_channel * c, struct sk_buff * skb);
```

## Arguments

*c*

The channel to use

*skb*

The packet to kick down the channel

## Description

Queue a packet for transmission. Because we have rather hard to hit interrupt latencies for the Z85230 per packet even in DMA mode we do the flip to DMA buffer if needed here not in the IRQ.

## z8530\_get\_stats

**Name** `z8530_get_stats` — Get network statistics

## Synopsis

```
struct net_device_stats * z8530_get_stats (struct z8530_channel * c);
```

## Arguments

*c*

The channel to use

## Description

Get the statistics block. We keep the statistics in software as the chip doesn't do it for us.

